

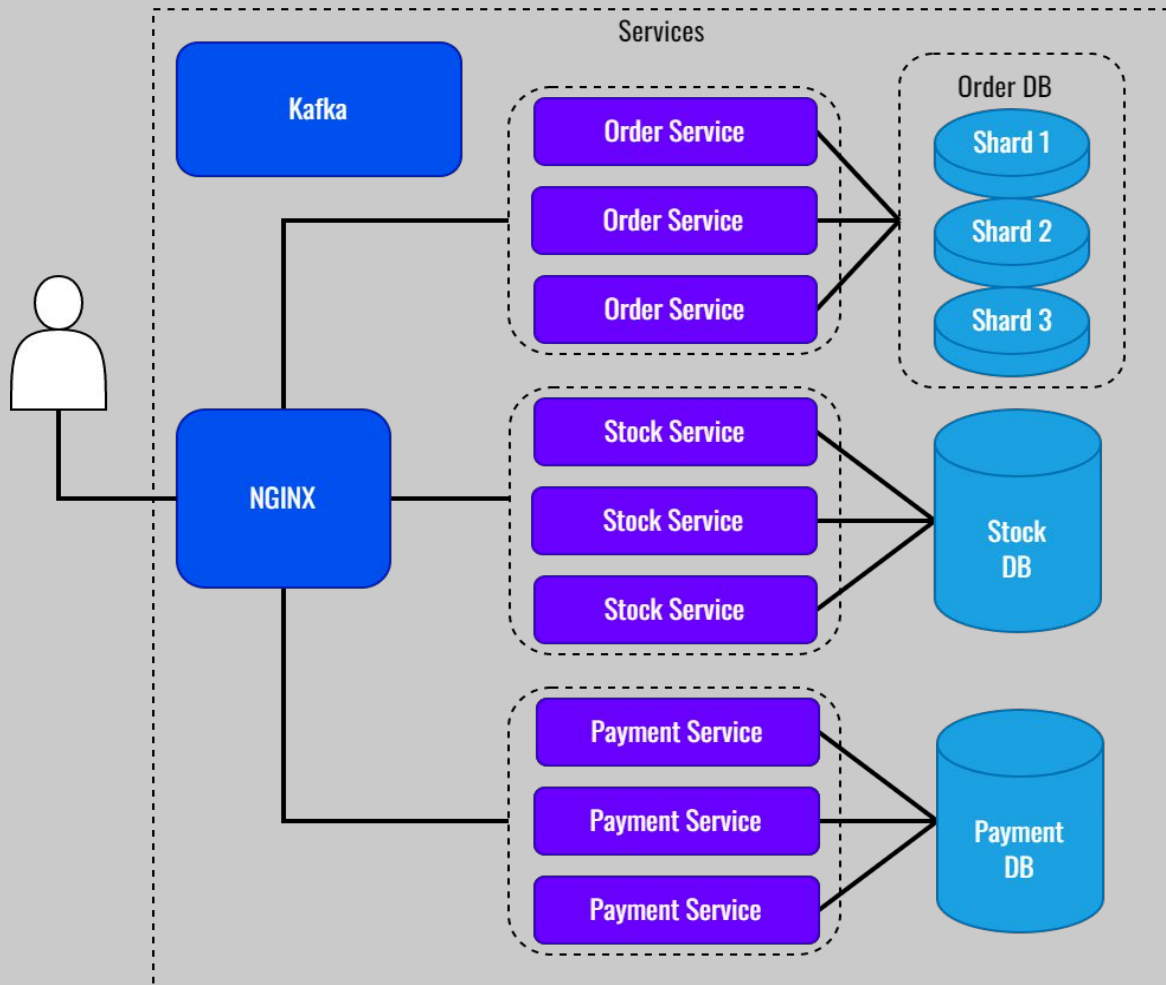
Web-Scale Data Management Project

...

Authors:
Alex Jeleniewski
Andrei Stefan
Paula Iacoban
Stef Kaptein
Tiago Gomes

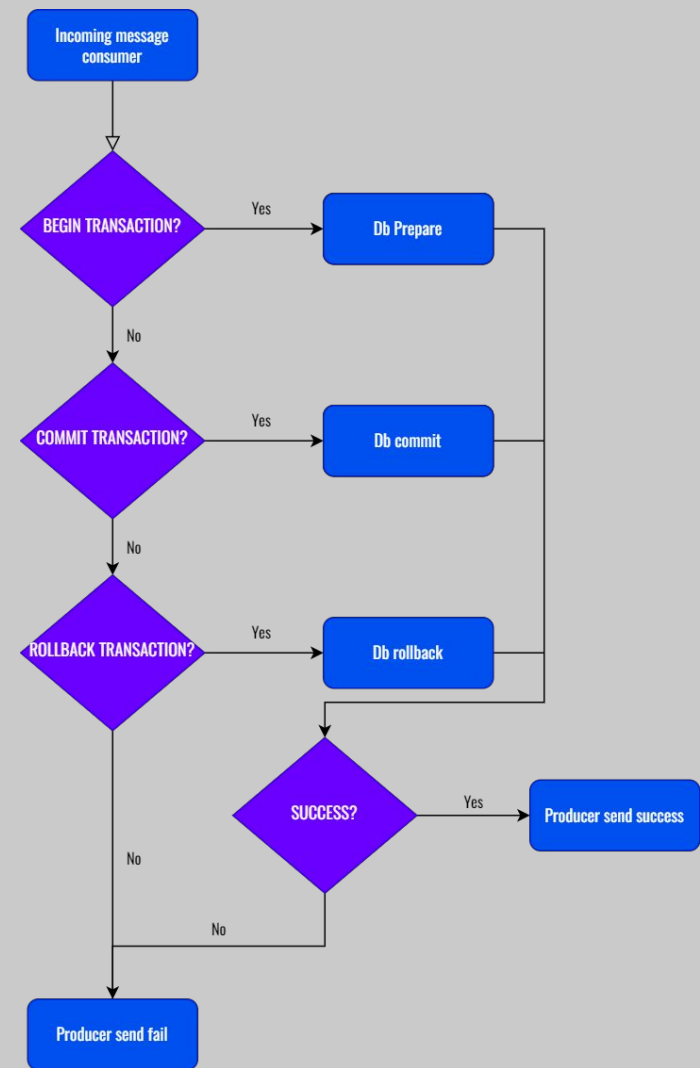
Design Architecture

- PostgreSQL
- Python Flask
- Kafka
- NGINX



Transaction Execution - 2PC

- Order service acts as coordinator -> Produce
- Stock and Payment consume instructions
- 3 instructions:
 - Prepare
 - Commit
 - Rollback
- 3 topics:
 - 1 for stock instructions
 - 1 for payment instructions
 - 1 for results




Transactional Design and Consistency Guarantees

- No illegal database state allowed
- 2 Phase Commit
- ACID Transactions
- Timeout detection

Scalability Methods

- Sharding of the Order database
- Multiple instances of all services
 - Actions can be performed by an arbitrary instance
 - Using partitions in Kafka
 - Example order service:

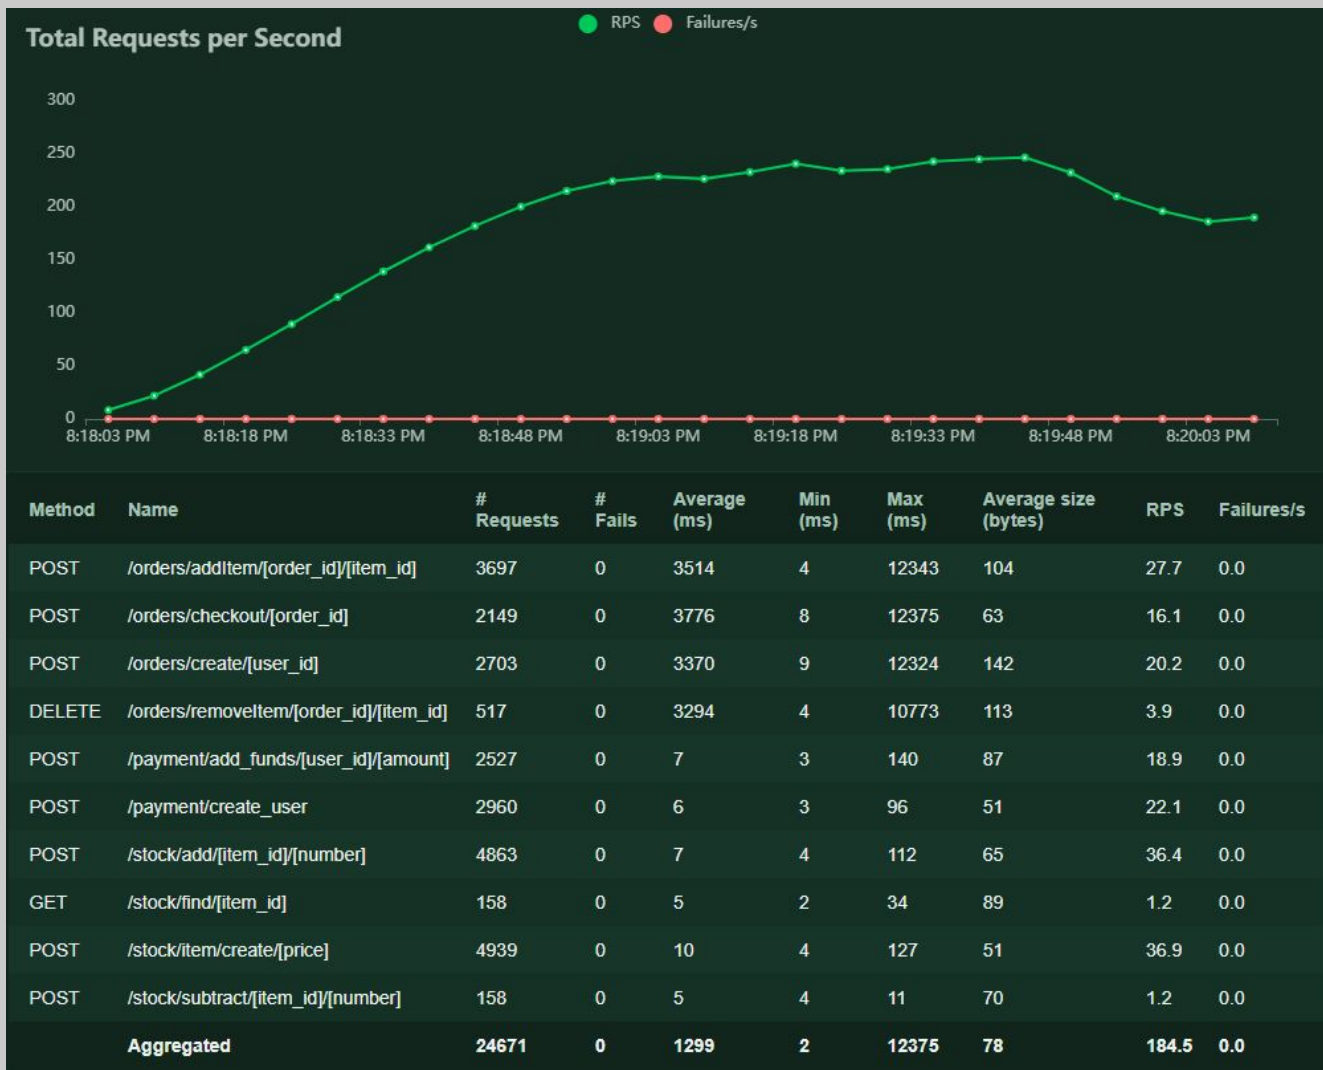


Fault-Tolerance Methods

- Two-Phase Commit
- Kafka's native capabilities
- Shards
- Multiple servers/service

Results & Measurements

- Balances out at around 200 requests per second



Results & Measurements

- Response time keeps increasing with number of users, but just linearly
- Order is the slowest

Response Times (ms)



Method	Name	50%ile (ms)	60%ile (ms)	70%ile (ms)	80%ile (ms)	90%ile (ms)	95%ile (ms)	99%ile (ms)	100%ile (ms)
POST	/orders/addItem/[order_id]/[item_id]	3100	4500	5300	6200	7200	8600	10000	12000
POST	/orders/checkout/[order_id]	3500	4700	5500	6500	7600	9200	11000	12000
POST	/orders/create/[user_id]	2800	4400	5100	6100	6900	7900	10000	12000
DELETE	/orders/removeItem/[order_id]/[item_id]	2600	4300	5200	6100	7300	8300	10000	11000
POST	/payment/add_funds/[user_id]/[amount]	6	7	7	8	10	12	18	140
POST	/payment/create_user	6	7	7	8	9	11	17	97
POST	/stock/add/[item_id]/[number]	6	7	7	8	10	11	20	110
GET	/stock/find/[item_id]	5	5	5	6	9	13	32	35
POST	/stock/item/create/[price]	7	8	10	16	22	27	35	130
POST	/stock/subtract/[item_id]/[number]	5	6	6	7	8	10	11	12
Aggregated		8	14	340	2500	5500	6800	9800	12000

Results & Measurements

- Consistent state in the database after each transaction, even if users went up to 1000



Limitations and future work

Limitations

- Tried sharding Payment, but didn't improve performance
- Sharding Stock would need distributed transactions
 - And might not improve performance

Future work

- Dynamic creation of shards (instead of fixed number)
- Dynamic scaling for the number of instances of services
- Sharding Payment and Stock for the cases where the application is under heavy load